

Milestone 2: Scientific Contextualization and Tradeoff Analysis of the Deferrable Server

Md Rafiul Islam

Electronic Engineering

Hamm-Lippstadt University of Applied Sciences

Germany

Abstract—This milestone document contextualizes the Deferrable Server within real-time scheduling. The focus is not polished final-paper prose, but structured scientific reasoning: neighboring approaches, assumptions, the central budget-preservation mechanism and its schedulability cost, engineering tradeoffs, embedded-system consequences, and open questions for feedback.

Index Terms—Deferrable Server, real-time scheduling, aperiodic tasks, periodic tasks, embedded systems, schedulability, tradeoff analysis

I. MILESTONE 2: SCIENTIFIC CONTEXTUALIZATION AND TRADEOFF ANALYSIS

A. Purpose and Scope

This milestone places the Deferrable Server (DS) inside the broader landscape of real-time scheduling and analyzes the tradeoffs it introduces, rather than collecting loosely related papers.

Core problem. On a uniprocessor, serve unpredictable *aperiodic* requests with low response time *without* causing the co-scheduled *periodic* hard real-time tasks to miss deadlines.

Scope and model.

- Single preemptive processor; fixed-priority baseline (Rate/Deadline Monotonic [1]), with Earliest Deadline First (EDF) noted as the dynamic-priority counterpart.
- Periodic task τ_i : known $(C_i, T_i, D_i) = \text{WCET, period, deadline}$. Aperiodic job: unknown release time, soft or firm deadline.
- A *server* is a schedulable entity with budget (capacity) Q and period P_s that reserves bandwidth $U_s = Q/P_s$ for aperiodic work.

B. Design Space: Where the Deferrable Server Sits

The mechanisms differ mainly in *how reserved capacity is created, preserved, or replenished*:

- **Idle-time, no reservation:** Background Server.
- **Reserved budget, non-preserving:** Polling Server.
- **Reserved budget, budget-preserving (fixed priority):** Deferrable Server, Priority Exchange, Sporadic Server.
- **Slack-based (fixed priority):** Slack Stealing.
- **Bandwidth reservation under EDF (dynamic priority):** Total Bandwidth Server, Constant Bandwidth Server.

The DS is the *simplest budget-preserving fixed-priority server*. That niche — better responsiveness than polling, far simpler

than sporadic — is exactly what motivates studying it and its tradeoffs.

C. Neighboring Approaches

- **Background Server:** aperiodic runs only when the CPU is idle. Zero interference on periodic tasks, but response time is poor and effectively unbounded under high periodic load.
- **Polling Server (PS):** full budget Q at each period; if no request is pending at the polling instant, the budget is *discarded*. Analyzable as an ordinary periodic task, but a request arriving just after a poll waits up to P_s .
- **Deferrable Server (DS):** full budget Q at each period, but unused budget is *preserved* until the period end. Improves aperiodic response; breaks the periodic-task equivalence (see Sec. I-D) [2], [3].
- **Priority Exchange (PE):** preserves budget by trading high-priority capacity for execution at lower priority levels when no aperiodic work is pending. Better utilization bound than DS, more bookkeeping, usually worse response than DS [2].
- **Sporadic Server (SS):** replenishes only the *consumed* amount, one period after consumption starts. Avoids the double-hit, so it behaves like a periodic task and *keeps the full Liu-Layland bound*; needs a replenishment queue [4].
- **Slack Stealing:** services aperiodic jobs at top priority using slack computed from the periodic schedule. Optimal aperiodic response under fixed priority, but high memory/compute cost [5].
- **Total Bandwidth Server (TBS, EDF):** assigns each job a deadline $d_k = \max(r_k, d_{k-1}) + C_k/U_s$ and schedules under EDF. Cheap, high utilization [6].
- **Constant Bandwidth Server (CBS, EDF):** enforces a (Q, T) reservation by postponing the deadline on overrun, giving temporal isolation; basis of Linux SCHED_DEADLINE [7].

D. Central Mechanism and Its Cost: Budget Preservation and the Double-Hit

Mechanism. Unlike the PS, the DS keeps unused budget within its period instead of discarding it. This is the entire source of its improved responsiveness.

TABLE I
SERVER-BASED MECHANISMS FOR APERIODIC SERVICE. THE DEFERRABLE SERVER IS THE SIMPLEST BUDGET-PRESERVING FIXED-PRIORITY OPTION.

Approach	Budget / replenishment	Schedulability basis	Aperiodic response	Cost / complexity
Background	none; runs on idle CPU	no extra analysis	poor, unbounded	trivial
Polling (PS)	full Q each P_s ; discarded if idle at poll	equivalent periodic task	poor if request misses poll	low
Deferrable (DS)	full Q each P_s ; <i>preserved</i> within period	lowered RM bound (double-hit)	good	low
Priority Exchange	high-prio budget exchanged to lower levels	better bound than DS	medium	medium
Sporadic (SS)	only consumed amount, one P_s after use	full Liu–Layland bound	good	medium (replenish. queue)
Slack Stealing	none; uses computed slack	optimal (fixed prio.)	best	high (slack tables)
TBS (EDF)	deadline $d_k = \max(r_k, d_{k-1}) + C_k/U_s$	EDF, $U_{\text{tot}} \leq 1$	good	low–medium
CBS (EDF)	(Q, T) reservation, deadline postponed	EDF, isolation	good	medium

The cost (back-to-back / “double-hit”). Preserved budget can be spent at the very *end* of period k ; the budget is then reset to a full Q at the *start* of period $k+1$ and can be spent immediately. The result is up to $2Q$ of *contiguous* highest-priority execution in a near-zero window (Fig. 1). A periodic task with identical (Q, P_s) cannot produce this burst.

Consequence for schedulability. Because of the double-hit, a DS *cannot* be modeled as an equivalent periodic task, and the schedulable periodic utilization drops. For a highest-priority DS with P_s no larger than all task periods, the least upper bound on schedulable periodic utilization is [2], [3]

$$U_p^{\text{lub}} = \ln\left(\frac{U_s + 2}{2U_s + 1}\right), \quad U_s = \frac{Q}{P_s}. \quad (1)$$

- $U_s \rightarrow 0 \Rightarrow U_p^{\text{lub}} \rightarrow \ln 2 \approx 0.693$ (the classic Liu–Layland bound is recovered).
- $U_s = 0.5 \Rightarrow U_p^{\text{lub}} = \ln 1.25 \approx 0.223$: the periodic set is limited to $\sim 22\%$ utilization.
- Eq. (1) is monotonically decreasing in U_s : more aperiodic bandwidth $\Rightarrow$ less guaranteed periodic bandwidth.

Why SS is “more controlled.” The Sporadic Server schedules its replenishment one period after consumption *begins*, which structurally prevents the $2Q$ burst. That is the precise technical reason SS preserves the full Liu–Layland bound while DS does not [4].

E. Assumptions (and Which Are Fragile in Practice)

- Single preemptive processor. [*Fragile: multicore breaks the single-resource analysis.*]
- Server-management and context-switch overhead negligible. [*Fragile on small MCUs.*]
- Periodic WCETs known and tight. [*Fragile: WCET estimation is hard on cached/pipelined cores.*]




Fig. 1. Deferrable Server “double-hit”: unused budget pushed to the end of one period, plus the full replenishment at the start of the next, yields up to $2Q$ of contiguous highest-priority execution. A same-parameter periodic task cannot generate this, which is why the DS lowers the schedulable periodic utilization (Eq. (1)).

- Aperiodic demand bounded by the server; excess simply waits. [*Firm/hard aperiodic deadlines are not guaranteed.*]
- Independent tasks, no shared resources / blocking. [*Fragile: real systems need PIP/PCP for mutexes.*]
- Negligible release jitter; server period $\leq$ task periods for Eq. (1). [*Often only approximately true.*]

F. Consequences for Embedded and Real-Time Systems

- **Memory footprint.** DS needs only a budget counter and a period timer ($O(1)$); SS needs a bounded replenishment queue; Slack Stealing needs slack tables — the heaviest.
- **Runtime overhead.** Roughly Background $<$ Polling $\leq$ DS $<$ SS $<$ Slack Stealing. The scheduler’s own WCET matters on constrained MCUs.
- **Predictability / certification.** Safety standards (e.g. DO-178C, ISO 26262) favour fully analyzable mechanisms; if DS is used, its worst-case $2Q$ interference must be explicitly budgeted.
- **RTOS mapping.** POSIX SCHED_SPORADIC $\approx$ Sporadic Server; Linux SCHED_DEADLINE $\approx$ CBS; the DS

TABLE II
ENGINEERING TRADEOFFS EXPOSED BY THE DEFERRABLE SERVER AND ITS NEIGHBOURS.

Tradeoff axis	What you gain (cheaper / simpler end)	What you pay (controlled / costly end)
Aperiodic response vs. periodic schedulability	DS preserves budget $\Rightarrow$ fast aperiodic service	lower guaranteed periodic utilization via Eq. (1)
Simplicity vs. no schedulability penalty	DS: trivial period-boundary replenishment	SS: replenishment bookkeeping, but full Liu-Layland bound
Static vs. dynamic priority	RM + DS: simple, certifiable, widely supported	EDF + TBS/CBS: higher utilization, isolation, more complex runtime
Offline guarantee vs. runtime flexibility	large Q : better aperiodic latency at runtime	large Q raises U_s and shrinks the periodic bound
Memory/overhead vs. optimal response	DS: $O(1)$ state (counter + timer)	Slack Stealing: optimal response but slack tables / computation

itself is rarely a native primitive and is usually *emulated*.

- **Energy / event handling.** A budgeted server consolidates bursty aperiodic work into predictable windows, which is friendlier to DVFS/sleep policies than ad-hoc interrupt handling.

Verdict. The DS is a good fit when (a) the platform uses fixed-priority scheduling, (b) soft aperiodic latency matters, and (c) U_s is small enough that the Eq. (1) penalty is tolerable. It is a poor *sole* mechanism for hard or safety-critical aperiodic paths unless the worst-case $2Q$ interference is fully accounted for; there SS or a reservation (CBS) is usually preferable.

G. Related-Work Status and Included vs. Excluded References

Status. The narrow core (the Lehoczky-Sha-Strosnider server family, synthesized by Buttazzo [8]) is *mature and nearly closed*: the design space is well characterized. “Sparse” here means few *new primitives*, not few papers — which is why this milestone compares a fixed, well-defined set instead of accumulating weakly related work.

Included (and why).

- DS papers [2], [3] — definitional; give the bound and the double-hit.
- Sporadic Server [4] — the direct “fix” to the DS schedulability penalty.
- Liu-Layland [1] — the bound the DS is measured against.
- TBS / CBS / Slack Stealing [5]–[7] — neighbouring families that bound the design space (EDF side, optimal response).
- Buttazzo [8] — unified analysis and notation.

Excluded (and why). General-purpose OS schedulers (e.g. Linux CFS), cloud/batch schedulers, global/partitioned multiprocessor scheduling, and learning-based schedulers — different assumptions and goals (throughput/fairness, not hard deadlines on one CPU).

H. Open Questions for Feedback

- For a concrete task set, is there a closed form for choosing (Q, P_s) that meets both Eq. (1) and an aperiodic-latency target, or must it be found by search?

TABLE III
MAPPING OF M2 CHECKLIST ITEMS TO THIS DOCUMENT.

M2 checklist item	Addressed in
Compared approaches/assumptions (not just summarized)	Tab. I, II
≥ 1 neighbouring approach / paradigm	Neighboring Approaches
Related-work strength (strong/weak/sparse) and why	Related-Work Status
Scientific / engineering tradeoffs	Tab. II, Eq. (1)
Embedded constraints (timing, memory, energy)	Consequences
Why selected references are relevant	Included references
Why excluded references are less relevant	Excluded references

- Can the extra deadline-miss risk caused specifically by the DS’s deferred budget (relative to a PS) be bounded per task?
- On a small MCU, when is the DS schedulability penalty cheaper than paying for SS replenishment bookkeeping?
- How does adding resource sharing (PIP/PCP blocking) interact with the server’s budget accounting?
- Does the double-hit penalty change under deadline-monotonic scheduling with $D_i < T_i$?

REFERENCES

- [1] C. L. Liu and J. W. Layland, “Scheduling algorithms for multiprogramming in a hard-real-time environment,” *J. ACM*, vol. 20, no. 1, pp. 46–61, 1973.
- [2] J. P. Lehoczky, L. Sha, and J. K. Strosnider, “Enhanced aperiodic responsiveness in hard real-time environments,” in *Proc. IEEE Real-Time Syst. Symp. (RTSS)*, 1987, pp. 261–270.
- [3] J. K. Strosnider, J. P. Lehoczky, and L. Sha, “The deferrable server algorithm for enhanced aperiodic responsiveness in hard real-time environments,” *IEEE Trans. Comput.*, vol. 44, no. 1, pp. 73–91, 1995.
- [4] B. Sprunt, L. Sha, and J. P. Lehoczky, “Aperiodic task scheduling for hard-real-time systems,” *Real-Time Syst.*, vol. 1, no. 1, pp. 27–60, 1989.
- [5] J. P. Lehoczky and S. Ramos-Thuel, “An optimal algorithm for scheduling soft-aperiodic tasks in fixed-priority preemptive systems,” in *Proc. IEEE Real-Time Syst. Symp. (RTSS)*, 1992, pp. 110–123.
- [6] M. Spuri and G. C. Buttazzo, “Scheduling aperiodic tasks in dynamic priority systems,” *Real-Time Syst.*, vol. 10, no. 2, pp. 179–210, 1996.

- [7] L. Abeni and G. Buttazzo, "Integrating multimedia applications in hard real-time systems," in *Proc. IEEE Real-Time Syst. Symp. (RTSS)*, 1998, pp. 4–13.
- [8] G. C. Buttazzo, *Hard Real-Time Computing Systems: Predictable Scheduling Algorithms and Applications*, 3rd ed. Springer, 2011.